

DSC630FinalProject

May 29, 2025

```
[52]: #DSC630 Final Project
# Armin Heldovac

# Import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix, classification_report

# Load dataset
df = pd.read_csv('C:\Users\armin\Downloads\surveylungcancer.csv')
df.head()
```

```
[52]:   GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE \
0      M    69          1                  2          2                  1
1      M    74          2                  1          1                  1
2      F    59          1                  1          1                  2
3      M    63          2                  2          2                  1
4      F    63          1                  2          1                  1

      CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING \
0              1        2        1        2                  2                  2
1              2        2        2        1                  1                  1
2              1        2        1        2                  1                  2
3              1        1        1        1                  2                  1
4              1        1        1        2                  1                  2

  SHORTNESS OF BREATH  SWALLOWING DIFFICULTY  CHEST PAIN  LUNG_CANCER
0                      2                  2          2          YES
1                      2                  2          2          YES
2                      2                  1          2          NO
```

3	1	2	2	NO
4	2	1	1	NO

```
[54]: # Remove duplicates
df = df.drop_duplicates()

# Check and handle missing values (example: remove rows with missing values)
df = df.dropna()

# Encode categorical variables
label_encoders = {}
for column in ['GENDER', 'SMOKING', 'ALCOHOL CONSUMING', 'COUGHING', 'SHORTNESS_OF_BREATH', 'CHEST_PAIN', 'LUNG_CANCER']:
    le = LabelEncoder()
    df[column] = le.fit_transform(df[column])
    label_encoders[column] = le

# Feature scaling for numerical variables
scaler = StandardScaler()
df['AGE'] = scaler.fit_transform(df[['AGE']])
```


```
[56]: #Data Visualization

# Class distribution
sns.countplot(x='LUNG_CANCER', data=df)
plt.title('Distribution of Lung Cancer Diagnosis')
plt.xlabel('Lung Cancer (0=No, 1=Yes)')
plt.ylabel('Number of Patients')
plt.show()


# Age distribution by diagnosis
sns.histplot(data=df, x='AGE', hue='LUNG_CANCER', kde=True, multiple='stack')
plt.title('Age Distribution by Lung Cancer Diagnosis')
plt.xlabel('Standardized Age')
plt.ylabel('Number of Patients')
plt.show()


# Correlation heatmap
plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Feature Correlation Heatmap')
plt.show()
```

Distribution of Lung Cancer Diagnosis



Age Distribution by Lung Cancer Diagnosis





[57]: #Model Building and Evaluation

```
# Define features and target
X = df.drop('LUNG_CANCER', axis=1)
y = df['LUNG_CANCER']

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train logistic regression model
lr_model = LogisticRegression(class_weight='balanced', random_state=42)
lr_model.fit(X_train, y_train)


# Make predictions
y_pred = lr_model.predict(X_test)
```

```
y_prob = lr_model.predict_proba(X_test)[:,1]
```


```
[58]: # Evaluate model
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC-AUC Score:", roc_auc_score(y_test, y_prob))
```

Accuracy: 0.8928571428571429
Precision: 1.0
Recall: 0.8636363636363636
F1 Score: 0.926829268292683
ROC-AUC Score: 0.9791666666666667

```
[62]: # Confusion matrix visualization
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=['No', 'Yes'],
            yticklabels=['No', 'Yes'])
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```



```
[64]: # Feature importance
feature_importance = pd.Series(lr_model.coef_[0], index=X.columns).
    ↪sort_values(ascending=False)
feature_importance.plot(kind='bar')
plt.title('Feature Importance (Logistic Regression Coefficients)')
plt.show()
```



```
[66]: # Classification report
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:					
	precision	recall	f1-score	support	
0	0.67	1.00	0.80	12	
1	1.00	0.86	0.93	44	
accuracy			0.89	56	
macro avg	0.83	0.93	0.86	56	
weighted avg	0.93	0.89	0.90	56	

```
[68]: from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_curve, auc
from sklearn.model_selection import cross_val_score, GridSearchCV

# Train Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_preds = rf_model.predict(X_test)
rf_probs = rf_model.predict_proba(X_test)[:, 1]

# Train KNN
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train, y_train)
knn_preds = knn_model.predict(X_test)
knn_probs = knn_model.predict_proba(X_test)[:, 1]

# Evaluate
print("Random Forest:")
print("Accuracy:", accuracy_score(y_test, rf_preds))
print("Precision:", precision_score(y_test, rf_preds))
print("Recall:", recall_score(y_test, rf_preds))
print("F1 Score:", f1_score(y_test, rf_preds))
print("ROC-AUC:", roc_auc_score(y_test, rf_probs))


print("\nK-Nearest Neighbors:")
print("Accuracy:", accuracy_score(y_test, knn_preds))
print("Precision:", precision_score(y_test, knn_preds))
print("Recall:", recall_score(y_test, knn_preds))
print("F1 Score:", f1_score(y_test, knn_preds))
print("ROC-AUC:", roc_auc_score(y_test, knn_probs))
```

Random Forest:
Accuracy: 0.8571428571428571
Precision: 0.8461538461538461

```
Recall: 1.0
F1 Score: 0.9166666666666666
ROC-AUC: 0.9820075757575757
```

```
K-Nearest Neighbors:
Accuracy: 0.8035714285714286
Precision: 0.8113207547169812
Recall: 0.9772727272727273
F1 Score: 0.8865979381443299
ROC-AUC: 0.9176136363636364
```

```
[70]: # ROC Curves
fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_model.predict_proba(X_test)[:,1])
fpr_rf, tpr_rf, _ = roc_curve(y_test, rf_probs)
fpr_knn, tpr_knn, _ = roc_curve(y_test, knn_probs)
plt.figure(figsize=(10,6))
plt.plot(fpr_lr, tpr_lr, label='LogReg (AUC={:.2f})'.
         format(roc_auc_score(y_test, lr_model.predict_proba(X_test)[:,1])))
plt.plot(fpr_rf, tpr_rf, label='Random Forest (AUC={:.2f})'.
         format(roc_auc_score(y_test, rf_probs)))
plt.plot(fpr_knn, tpr_knn, label='KNN (AUC={:.2f})'.
         format(roc_auc_score(y_test, knn_probs)))
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.grid()
plt.show()
```



```
[72]: # Cross-validation Recall
print("Logistic Regression CV Recall:", cross_val_score(lr_model, X, y, cv=5,
    scoring='recall').mean())
print("Random Forest CV Recall:", cross_val_score(rf_model, X, y, cv=5,
    scoring='recall').mean())
print("KNN CV Recall:", cross_val_score(knn_model, X, y, cv=5,
    scoring='recall').mean())

# Grid Search for RF
param_grid_rf = {'n_estimators': [50, 100], 'max_depth': [None, 10],
    'min_samples_split': [2, 5]}
grid_rf = GridSearchCV(RandomForestClassifier(random_state=42), param_grid_rf,
    cv=5, scoring='recall')
grid_rf.fit(X_train, y_train)
print("Best RF Params:", grid_rf.best_params_)
print("Best RF CV Recall:", grid_rf.best_score_)

# Grid Search for KNN
param_grid_knn = {'n_neighbors': [3, 5, 7]}
grid_knn = GridSearchCV(KNeighborsClassifier(), param_grid_knn, cv=5,
    scoring='recall')
grid_knn.fit(X_train, y_train)
print("Best KNN Params:", grid_knn.best_params_)
```

```
print("Best KNN CV Recall:", grid_knn.best_score_)
```

Logistic Regression CV Recall: 0.8780141843971631
Random Forest CV Recall: 0.9454787234042554
KNN CV Recall: 0.962145390070922
Best RF Params: {'max_depth': None, 'min_samples_split': 5, 'n_estimators': 100}
Best RF CV Recall: 0.9690958164642375
Best KNN Params: {'n_neighbors': 7}
Best KNN CV Recall: 0.9794871794871796

[]:

[]: